

# ECE521 W17 Tutorial 4

Min Bai & Kaustav Kundu



UNIVERSITY OF  
**TORONTO**

*\*Reference: CSC411 lecture slides, Stanford's CS321n course*

# Agenda

- Logistic Regression
- Neural Networks
- Assignment 1 Questions


# Logistic Regression

- We assumed a particular functional form: sigmoid applied to a linear function of the data

$$y(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$$

where the sigmoid is defined as

$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$



- The output is a smooth function of the inputs and the weights

# Logistic Regression

- We assumed a particular functional form: sigmoid applied to a linear function of the data

$$y(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$$

where the sigmoid is defined as

$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$

- ▶ One parameter per data dimension (feature)
- ▶ Features can be discrete or continuous
- ▶ Output of the model: value  $y \in [0, 1]$

# Logistic Regression

- If we have a value between 0 and 1, let's use it to model class probability

$$p(C = 0|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0) \quad \text{with} \quad \sigma(z) = \frac{1}{1 + \exp(-z)}$$

- Substituting we have

$$p(C = 0|\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x} - w_0)}$$

- Suppose we have two classes, how can I compute  $p(C = 1|\mathbf{x})$ ?
- Use the marginalization property of probability


$$p(C = 1|\mathbf{x}) + p(C = 0|\mathbf{x}) = 1$$

- Thus



$$p(C = 1|\mathbf{x}) = 1 - \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x} - w_0)} = \frac{\exp(-\mathbf{w}^T \mathbf{x} - w_0)}{1 + \exp(-\mathbf{w}^T \mathbf{x} - w_0)}$$

# Logistic Regression

- What is the **decision boundary** for logistic regression?
- $p(C = 1|\mathbf{x}, \mathbf{w}) = p(C = 0|\mathbf{x}, \mathbf{w}) = 0.5$
- $p(C = 0|\mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0) = 0.5$ , where  $\sigma(z) = \frac{1}{1+\exp(-z)}$
- Decision boundary:  $\mathbf{w}^T \mathbf{x} + w_0 = 0$
- Logistic regression has a **linear decision boundary**



# Logistic Regression



If the right answer is 1 and the model says 1.5, it loses, so it changes the boundary to avoid being “too correct” (tilts away from outliers)

# Logistic Regression

- We can also look at

$$p(\mathbf{w}|\{t\}, \{\mathbf{x}\}) \propto p(\{t\}|\{\mathbf{x}\}, \mathbf{w}) p(\mathbf{w})$$

with  $\{t\} = (t^{(1)}, \dots, t^{(N)})$ , and  $\{\mathbf{x}\} = (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)})$

- We can define priors on parameters  $\mathbf{w}$
- This is a form of regularization
- Helps avoid large weights and [overfitting](#)

$$\max_{\mathbf{w}} \log \left[ p(\mathbf{w}) \prod_i p(t^{(i)}|\mathbf{x}^{(i)}, \mathbf{w}) \right]$$

- What's  $p(\mathbf{w})$ ?



# Logistic Regression

- For example, define prior: normal distribution, zero mean and identity covariance  $p(\mathbf{w}) = \mathcal{N}(0, \alpha^{-1}\mathbf{I})$

$$f(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{(\mu-x)^2}{2\sigma^2}}$$

- This prior pushes parameters towards zero (why is this a good idea?)
- Including this prior the new gradient is


$$w_j^{(t+1)} \leftarrow w_j^{(t)} - \lambda \frac{\partial \ell(\mathbf{w})}{\partial w_j} - \lambda \alpha w_j^{(t)}$$

where  $t$  here refers to iteration of the gradient descent

- The parameter  $\alpha$  is the importance of the regularization, and it's a **hyper-parameter**
- How do we decide the best value of  $\alpha$  (or a hyper-parameter in general)?

# Problems with Linear Classifiers


- Linear classifiers (e.g., logistic regression) classify inputs based on linear combinations of features  $x_i$
- Many decisions involve non-linear functions of the input
- Canonical example: do 2 input elements have the same value?



- The positive and negative cases cannot be separated by a plane
- What can we do?

# Intro to Neural Networks


- Neural networks are computational models *inspired* by the human brain.
- Our brain has  $\sim 10^{11}$  neurons, each of which is connected to  $\sim 10^4$  other neurons.




Cartoon diagram of a biological neuron

# Intro to Neural Networks

- Neural networks are functions of the neurons, also called hidden units.



Cartoon diagram of a biological neuron







Mathematical model of a biological neuron

# Intro to Neural Networks

- Neural networks are functions of the neurons, also called hidden units.
- Commonly used activation functions


- Sigmoid:  $\sigma(z) = \frac{1}{1 + \exp(-z)}$
- Tanh:  $\tanh(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}$
- ReLU (Rectified Linear Units):

$$\text{ReLU}(x) = \max(0, z)$$



# Multi layer Perceptrons

- Going deeper: a 3-layer neural network with two layers of hidden units




**Figure :** A 3-layer neural net with 3 input units, 4 hidden units in the first and second hidden layer and 1 output unit

- Naming conventions; a  $N$ -layer neural network:
  - ▶  $N - 1$  layers of hidden units
  - ▶ One output layer

# Multi layer Perceptrons

- Neural Network with **at least one hidden layer** is a universal approximator (can represent any function)<sup>1</sup>.



- Capacity of the network increases with more hidden units and more hidden layers
- Why go deeper? Read Jimmy's [paper](#) or the paper on [the loss surface of multilayer networks](#).


<sup>1</sup>Proof by Cibenko in [Approximation by Superpositions of a Sigmoidal Function](#). Intuitive Explanation from Michael Nielson: [link](#)

# Neural Networks

- We only need to know two algorithms
  - ▶ **Forward pass:** performs inference
  - ▶ **Backward pass:** performs learning



# Inference: Forward Pass



- Output of the network can be written as:

$$h_j(\mathbf{x}) = f(v_{j0} + \sum_{i=1}^D x_i v_{ji})$$

$$o_k(\mathbf{x}) = g(w_{k0} + \sum_{j=1}^J h_j(\mathbf{x}) w_{kj})$$


( $j$  indexing hidden units,  $k$  indexing the output units,  $D$  number of inputs)

- Activation functions  $f$ ,  $g$ : sigmoid/logistic, tanh, or rectified linear (ReLU)

$$\sigma(z) = \frac{1}{1 + \exp(-z)}, \quad \tanh(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}, \quad \text{ReLU}(z) = \max(0, z)$$

# Special Case

- What is a single layer (no hidden) network with a sigmoid act. function?




- Network:

$$o_k(\mathbf{x}) = \frac{1}{1 + \exp(-z_k)}$$
$$z_k = w_{k0} + \sum_{j=1}^J x_j w_{kj}$$

# Special Case

- What is a single layer (no hidden) network with a sigmoid act. function?



- Network:

$$o_k(\mathbf{x}) = \frac{1}{1 + \exp(-z_k)}$$
$$z_k = w_{k0} + \sum_{j=1}^J x_j w_{kj}$$

- Logistic regression!

# Training: Backward Pass

- Find weights:

$$\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} \sum_{n=1}^N \operatorname{loss}(\mathbf{o}^{(n)}, \mathbf{t}^{(n)})$$

where  $\mathbf{o} = f(\mathbf{x}; \mathbf{w})$  is the output of a neural network

- Define a loss function, eg:

- ▶ Squared loss:  $\sum_k \frac{1}{2} (o_k^{(n)} - t_k^{(n)})^2$
- ▶ Cross-entropy loss:  $-\sum_k t_k^{(n)} \log o_k^{(n)}$

- Gradient descent:


$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \frac{\partial E}{\partial \mathbf{w}^t}$$

where  $\eta$  is the learning rate (and  $E$  is error/loss)

# Training: Backward Pass

$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2$ ,  $y = 5$ ,  $z = -4$



# Training: Backward Pass


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



# Training: Backward Pass


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



# Training: Backward Pass


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2$ ,  $y = 5$ ,  $z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}$ ,  $\frac{\partial f}{\partial y}$ ,  $\frac{\partial f}{\partial z}$





# Training: Backward Pass


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial y}$$

# Training: Backward Pass


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2$ ,  $y = 5$ ,  $z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}$ ,  $\frac{\partial f}{\partial y}$ ,  $\frac{\partial f}{\partial z}$



Chain rule:

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial y}$$

$$\frac{\partial f}{\partial y}$$

# Training: Backward Pass


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2$ ,  $y = 5$ ,  $z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}$ ,  $\frac{\partial f}{\partial y}$ ,  $\frac{\partial f}{\partial z}$




Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$

$$\frac{\partial f}{\partial x}$$


# Training: Backward Pass



# Another Example


- Let's assume a two class problem

$$f(w, x) = \frac{1}{1 + e^{-(w_0 x_0 + w_1 x_1 + w_2)}}$$



# Another Example

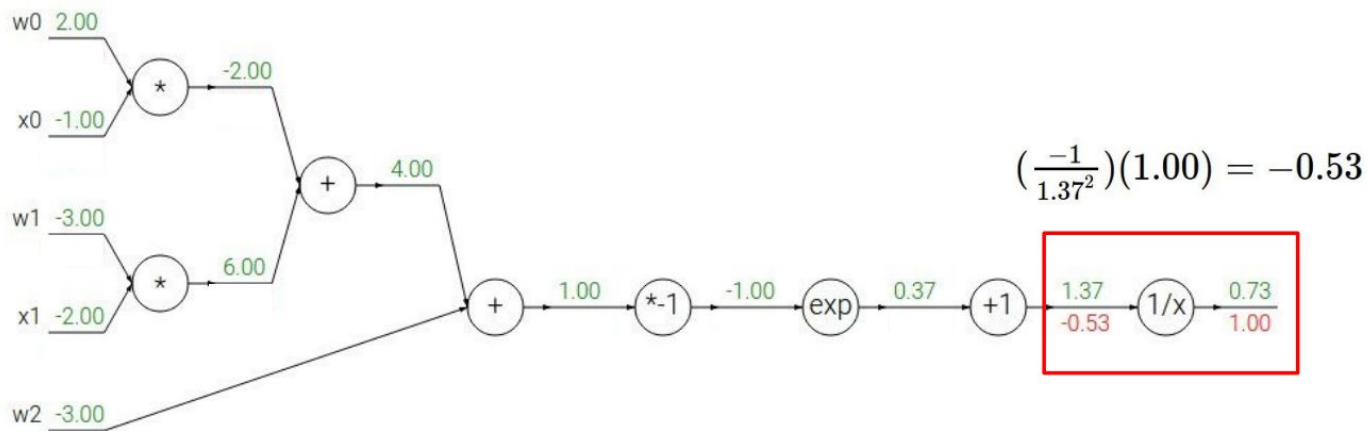
$$f(w, x) = \frac{1}{1 + e^{-(w_0 x_0 + w_1 x_1 + w_2)}}$$



$f(x) = e^x$	$\rightarrow$	$\frac{df}{dx} = e^x$		$f(x) = \frac{1}{x}$	$\rightarrow$	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	$\rightarrow$	$\frac{df}{dx} = a$		$f_c(x) = c + x$	$\rightarrow$	$\frac{df}{dx} = 1$

# Another Example

$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2)}}$$



$$f(x) = e^x \rightarrow \frac{df}{dx} = e^x$$


$$f_a(x) = ax \rightarrow \frac{df}{dx} = a$$

$$f(x) = \frac{1}{x} \rightarrow \frac{df}{dx} = -1/x^2$$

$$f_c(x) = c + x \rightarrow \frac{df}{dx} = 1$$

# Another Example

$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2)}}$$



$$(e^{-1})(-0.53) = -0.20$$

$$f(x) = e^x$$

$\rightarrow$

$$\frac{df}{dx} = e^x$$

$$f_a(x) = ax$$

$\rightarrow$

$$\frac{df}{dx} = a$$

$$f(x) = \frac{1}{x}$$

$\rightarrow$

$$\frac{df}{dx} = -1/x^2$$

$$f_c(x) = c + x$$


$\rightarrow$

$$\frac{df}{dx} = 1$$



# Another Example

$$f(w, x) = \frac{1}{1 + e^{-(w_0 x_0 + w_1 x_1 + w_2)}}$$



$f(x) = e^x$	$\rightarrow$	$\frac{df}{dx} = e^x$		$f(x) = \frac{1}{x}$	$\rightarrow$	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	$\rightarrow$	$\frac{df}{dx} = a$		$f_c(x) = c + x$	$\rightarrow$	$\frac{df}{dx} = 1$


# Another Example

$$f(w, x) = \frac{1}{1 + e^{-(w_0 x_0 + w_1 x_1 + w_2 x_2)}}$$

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

sigmoid function

$$\frac{d\sigma(x)}{dx} = \frac{e^{-x}}{(1 + e^{-x})^2} = \left( \frac{1 + e^{-x} - 1}{1 + e^{-x}} \right) \left( \frac{1}{1 + e^{-x}} \right) = (1 - \sigma(x)) \sigma(x)$$



## Other useful derivatives

name	function	derivative
Sigmoid	$\sigma(z) = \frac{1}{1+\exp(-z)}$	$\sigma(z) \cdot (1 - \sigma(z))$
Tanh	$\tanh(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}$	$1 / \cosh^2(z)$
ReLU	$\text{ReLU}(z) = \max(0, z)$	$\begin{cases} 1, & \text{if } z > 0 \\ 0, & \text{if } z \leq 0 \end{cases}$

# Cross Entropy

- Entropy of a distribution  $p$ ,  $H(p)$ , measures expected number of bits to represent data from this distribution using the optimal code from  $p$  distribution.

$$H(p) = - \sum_x p(x) \log p(x) = -E_{x \sim p(x)} [\log p(x)]$$

- Cross entropy of  $p$  w.r.t.  $q$  measures how many bits you require on average to represent data from  $p$  distribution using the optimal code from  $q$  distribution.

$$H(p, q) = - \sum_x p(x) \log q(x) = -E_{x \sim p(x)} [\log q(x)]$$

# Cross Entropy

- Defined between two distributions:

$$\begin{aligned} H(p, q) &= - \sum_x p(x) \log q(x) \\ &= - \sum_x p(x) \log p(x) + \sum_x p(x) \log \frac{p(x)}{q(x)} \\ &= H(p) + D_{KL}(p||q) \end{aligned}$$

- $D_{KL}(p||q)$  is the extra number of bits required to code data from  $p$  distribution using the optimal code for  $q$  distribution.

# Cross Entropy

- Defined between two distributions:

$$\begin{aligned} H(p, q) &= - \sum_x p(x) \log q(x) \\ &= H(p) + D_{KL}(p||q) \end{aligned}$$

- $p = [0, \dots, 0, 1, 0, \dots, 0]$ , puts the entire probability mass on the ground truth class.
- $q(x_i) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$  is the estimated class "probabilities"

$$L(y, y_{gt}; w) = - \sum_i \sum_j y_{gt}^{(i)}(x) \log y^{(i)}(x) = - \sum_i \log(y_{y_{gt}^{(i)}}^{(i)})$$

# Cross Entropy + Weight Decay

- The total loss:

$$L(y, y_{\text{gt}}; w) = - \sum_i \log(y_{y_{\text{gt}}^{(i)}}^{(i)}) + ||w||^2$$

- Can be interpreted as a MAP estimation with a Gaussian prior on the weights.

# Weight Initialization

- Logistic Regression
  - Convex optimization - initialization doesn't matter much
  - Initialize weights and biases to zero
- Neural Network
  - Initialize to random numbers - symmetry breaking
    - Otherwise, all branches have same weight derivatives and updates
  - Sample from Gaussian / truncated Gaussian
  - Mean = 0
  - Variance?
    - Xavier initialization guideline:

$$\text{var}(w) = \sqrt{\frac{2}{n_{\text{in}} + n_{\text{out}}}}$$



# Data Normalization

- Want each dimension of input data to be zero centered and have approximately same variance
  - Eg  $N(0,1)$
  - For each channel:
    - Mean subtraction
    - Scaling

